Due data: 10/7/2020, end of the day.

> **For question 1-3 , please submit a PDF file via Canvas.**
> **For question 4 (programming question), please submit an .ipynb file via Canvas.**

Please answer the following questions:

1. [*6 points*] Prove Bayes' Theorem. Briefly explain why it is useful for machine learning problems, i.e., by converting posterior probability to likelihood and prior probability.

2. [10 points] In Lecture 2-2, we gave the normal equation (i.e., closed-form solution) for linear regression using MSE as the cost function. **Prove that the closed-form solution for Ridge Regression is $w = (\lambda I + X^T \cdot X)^{-1} \cdot X^T \cdot y$**, where I is the identity matrix, $X = (x^{(1)}, x^{(2)}, ..., x^{(m)})^T$ is the input data matrix, $x^{(i)} = (1, x_1, x_2, ..., x_n)$ is the ith data sample, and $y = (y^{(1)}, y^{(2)}, ..., y^m)$. Assume the hypothesis function $h_w(x) = w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_n x_n$, and $y^{(j)}$ is the measurement of $h_w(x)$ for the j th training sample. The cost function of the Ridge Regression is $E(w) = \sum_{i=1}^{m}(w^T \cdot x^{(i)} - y^{(i)})^2 + \lambda \sum_{i=1}^{m} w_i^2$.

3. [10 points] Recall the multi-class Softmax Regression model on page 16 of Lecture 3-2. Assume we have K different classes. The posterior probability is $\hat{p}_k = \delta(s_k(x))_k = \frac{\exp(s_k(x))}{\sum_{j=1}^{K} \exp(s_j(x))}$ for $k = 1, 2, ..., K$, where $s_k(x) = \theta_k^T \cdot x$, and input x is an n-dimension vector.
   1) To learn this Softmax Regression model, how many parameters we need to estimate? What are these parameters?
   2) Consider the cross-entropy cost function $J(\Theta)$ (see page 16 of Lecture 3-2) of m training samples $\{(x_i, y_i)\}_{i=1,2,...,m}$. Derive the gradient of $J(\Theta)$ regarding to θ_k as shown in page 17 of Lecture 3-2

Programming Problem:
4. [44 *points*] In this problem, we write a program to find the coefficients for a linear regression model for the dataset provided (data2.txt). Assume a linear model: y = w$_0$ + w$_1$*x. You need to
   1) Plot the data (i.e., x-axis for 1st column, y-axis for 2nd column),

   and use Python to implement the following methods to find the coefficients:

   2) Normal equation, and
   3) Gradient Descent using **batch** AND **stochastic** modes respectively:
      a) Determine an appropriate termination condition (e.g., when cost function is less than a threshold, and/or after a given number of iterations).
      b) Print the cost function vs. iterations for each mode; compare and discuss batch and stochastic modes in terms of the accuracy and the speed of convergence.
      c) Choose a best learning rate. For example, you can plot cost function vs. learning rate to determine the best learning rate.

Please implement the algorithms by yoursef and **do NOT use the fit() function** of the library.